

SAI TEJASRI YERRAMSETTI

San Diego, CA | ysaitejasrigrad@gmail.com | 619-394-5167 | [linkedin.com/in/saitejasri](https://www.linkedin.com/in/saitejasri) | github.com/saitejasri1

Data Scientist with 2+ years of experience building scalable ML models, A/B testing strategies, and deploying AI-driven solutions. Specialized in Generative AI, predictive modeling, and optimizing end-to-end machine learning pipelines for real-world impact.

Work Experience

Data Science Intern

Zip Launchpad, San Diego, CA — Jan 2024 – June 2024

- Developed and deployed NLP models to analyze open-ended survey responses at scale, surfacing key insights that improved program satisfaction by 20%.
- Automated data preprocessing pipelines using PySpark and Airflow, reducing manual effort by 40% and enabling real-time analytics for program evaluation.
- Built predictive models using ARIMA and Prophet, forecasting session attendance trends, improving capacity planning, and reducing underutilization by 25%.
- Designed interactive dashboards leveraging Etrack survey data (2020-2024), providing real-time insights that guided leadership in data-driven decision-making.

Data Science Intern

Digital Innovation Lab — Jun 2023 – Aug 2023

- Developed an LLM-powered mentor-mentee chatbot using TogetherAI API and RAG, handling 10K+ user interactions, increasing engagement by 40%, and reducing response latency by 30%.
- Fine-tuned Mistral 7B for summarization, sentiment analysis, and keyword extraction, improving model accuracy by 18% and reducing compute costs by 20%.
- Built a real-time sentiment analysis pipeline using Spark and PyTorch, analyzing mentor-mentee interactions at scale, identifying 3 key engagement factors, and improving retention by 25%.
- Optimized LLM deployment on AWS, leveraging GPU-accelerated inference, reducing model latency by 25%, and cutting cloud costs by 15%.
- Implemented a continuous learning loop with cloud-stored embeddings, automating real-time model retraining, boosting chatbot response relevance by 15%, and scaling to 100K+ conversations.

Software Product Engineer

Temenos, Hyderabad, India — Aug 2021 – Aug 2023

- Resolved over 200 Jira tickets and debugged workflows, improving the stability of Transact and Infinity products.
- Architected robust front-end microservices using Visualizer with seamless integration of backend systems, enabling smooth handling of up to 100 concurrent transactions during peak usage periods at California Coast Credit Union.
- Designed scalable microservices for retail and wealth management systems, reducing deployment time by 30%.
- Streamlined CI/CD pipelines, reducing deployment errors by 30% and enhancing software delivery speed.
- Engineered API-driven solutions, collaborating with cross-functional teams to enhance customer experience, increasing satisfaction scores by 12% across multiple deployments.

Projects

Pixel to Plate: Deep Learning Based- Recipe Discovery and Recommendation

- Developed a deep learning-based recipe recommendation system integrating computer vision and NLP, achieving 96.38% accuracy in ingredient recognition and personalized recipe generation.
- Implemented an ensemble model to generate curated recipe recommendations from ingredient images, enhancing accuracy and user engagement.

Audio Feedback Processing Platform (Hackathon Project)

- Led a team of 4 to build an application enabling patients to submit audio feedback, integrating speech-to-text, sentiment analysis, keyword extraction, text summarization, and categorization.
- Developed a hospital-facing dashboard that reduced manual feedback analysis time by 70%, enabling faster response to patient concerns.
- Leveraged WebSockets and PyTorch models to accelerate feedback loops by 40%, improving the accuracy of sentiment monitoring and overall operational efficiency.

Education

San Diego State University, San Diego, CA [CGPA 3.8]

Aug 2023 – Present

Master's in Big Data Analytics

Technical Skills

Programming and Frameworks: Python, R, SQL, Rust, TensorFlow, PyTorch, scikit-learn, Keras

Data Visualization and Backend Development: Pandas, NumPy, Matplotlib, Seaborn, Tableau, Power BI, RESTful APIs, Flask, Postman, Microservices, MySQL, PostgreSQL, MongoDB, Spark, Hadoop

Cloud and Deployment: AWS (S3, EC2, SageMaker), Azure, GCP, Flask, Docker, Streamlit, Dash

Other: LLM Optimization, NLP, Debugging, Git, Database Management, Github, CI/CD Pipelines (Jenkins)

Certifications

- Oracle Cloud Infrastructure 2024 Generative AI Professional (1Z0-1127-24)
- Databricks Academy Accreditation - Generative AI Fundamentals